

Detecting Misinterpreted but Genuine Citations Through Automated Claim Verification

Tushar Kumar

Roll No.: MHM2026010

M.Tech IT (Specialization in Human Machine Interaction and Immersive Technology)

Department of Information Technology, Indian Institute of Information Technology, Allahabad

Prayagraj, Uttar Pradesh, India

Email: mhm2026010@iiita.ac.in

Abstract—The rapid proliferation of academic literature and the widespread integration of automated writing assistants have intensified a critical yet understudied vulnerability in scientific discourse: misinterpreted but genuine citations. In this failure mode, a cited scholarly publication genuinely exists and possesses flawless bibliographic metadata, yet the in-text assertion made by the citing author misrepresents, exaggerates, generalizes, or directly contradicts the cited publication's empirical findings. Traditional bibliographic audits and automated citation managers verify only surface-level presence—such as DOI resolution and catalog metadata—remaining blind to semantic infidelity between claims and cited evidence. This paper provides a comprehensive investigation into detecting misinterpreted genuine citations through automated, evidence-grounded claim verification. We formalize the verification lifecycle as an end-to-end Natural Language Processing (NLP) pipeline comprising citation-context extraction, atomic claim decomposition, hierarchical full-text evidence retrieval, fine-grained cross-document natural language inference (NLI), and calibrated decision synthesis. We systematically evaluate lexical retrieval, dense bi-encoders, cross-attention encoders, Retrieval-Augmented Generation (RAG), and Large Language Model (LLM) reasoning architectures, exposing their inductive biases and operational failure modes. Furthermore, we construct an analytical taxonomy of citation misinterpretation—spanning numerical distortions, causal inflation, methodological conflation, and scope extrapolation—and demonstrate why standard surface metrics fail. Finally, we analyze open challenges including multimodal reasoning across figures and tables, epistemic hedging, commercial paywalls, and citation error cascades, articulating a strategic roadmap toward trustworthy, explainable, and human-in-the-loop scientific auditing systems.

Index Terms—*Citation Verification, Claim Verification, Citation Entailment, Natural Language Inference, Evidence Retrieval, Scientific NLP, Large Language Models, Citation Integrity, Scholarly Information Retrieval.*

I. INTRODUCTION

The contemporary scholarly ecosystem is undergoing an unprecedented transformation characterized by an exponential growth in published literature, the globalization of scientific collaboration, and the pervasive integration of artificial intelligence (AI) and automated text generation systems into scholarly writing workflows [1], [2]. Millions of peer-reviewed manuscripts are indexed annually across computer science, biomedicine, physical sciences, and engineering, expanding the frontiers of knowledge while placing immense strain on traditional mechanisms of peer review, editorial oversight, and academic quality control. In parallel, the deployment of large language models (LLMs) and automated drafting assistants has significantly lowered the barrier to text generation, enabling researchers to produce complex academic prose at remarkable velocity [1]. However, this technological acceleration has introduced systemic vulnerabilities into scholarly discourse, among which citation integrity has emerged as one of the most critical and challenging concerns [2], [3].

Within the normative framework of scientific inquiry, citations serve as the foundational bedrock of epistemic accountability [3], [4]. Citations acknowledge intellectual provenance, anchor novel hypotheses within established empirical frameworks, provide reproducible grounding for methodological decisions, and credit prior discoveries. An academic assertion lacking evidentiary backing represents mere conjecture; an assertion anchored by an authoritative citation claims the weight of peer-reviewed verification. Consequently, the epistemic validity of an entire scientific manuscript frequently hinges on the fidelity of the links connecting its textual claims to the evidentiary foundation provided by its cited literature [4], [5].

A. The Illusion of Citation Existence: Authenticity vs. Epistemic Support

Historically, automated systems for academic auditing and citation management—such as reference managers, plagiarism detection engines, and publisher metadata parsers—have operated under a fundamental simplification: equating citation presence with citation validity [5], [6]. These conventional tools verify surface-level bibliographic criteria: Does the cited Digital Object Identifier (DOI) resolve to a valid record? Does the cited paper exist in Crossref, Scopus, or PubMed? Are the author names, publication year, volume, and pagination structurally sound? While detecting fabricated, non-existent, or 'phantom' citations is vital—particularly in an era where ungrounded generative AI models frequently hallucinate plausible-sounding references [1], [9]—verifying that a publication exists is fundamentally distinct from verifying that it supports the proposition for which it is cited [6], [7].

This divergence reveals a critical conceptual boundary separating two independent axes of scholarly integrity:

- 1) Bibliographic Authenticity: The cited source is a verified, authentic, cataloged publication possessing correct metadata.
- 2) Semantic Entailment (Citation Support): The empirical findings, theoretical derivations, or methodological outcomes of the cited source genuinely substantiate, justify, and entail the specific claim made in the citing text.

A citation can be entirely authentic, flawlessly formatted, and topical to the broad research area, while remaining completely disconnected from—or even diametrically opposed to—the specific assertion made by the citing author [7], [8].

B. The Problem of Misinterpreted but Genuine Citations

This paper addresses the nuanced and pervasive failure mode: misinterpreted but genuine citations. In this failure mode, authors cite authentic, highly reputable, and correctly identified scholarly publications, yet the citing sentence misrepresents, exaggerates, overgeneralizes, or contradicts the source document's actual contributions [7], [8]. Rather than being an outright fabrication, the citation acts as a veneer of academic legitimacy, masking unsupported claims beneath the prestige of genuine literature.

Misinterpretation manifests across diverse structural and semantic dimensions in academic writing [8], [10]. Authors may claim that an intervention was clinically proven effective when the cited trial yielded statistically non-significant primary endpoints; assert a causal relationship when the cited study explicitly highlighted observational correlations; extrapolate pilot results observed in a specialized in vitro or murine cohort to universal human populations; or cite a speculative hypothesis discussed in an introductory paragraph as if it were an empirically established finding [7], [11]. Furthermore, quantitative parameters—such as effect sizes, odds ratios, model parameter counts, and error rates—are frequently misquoted, inverted, or stripped of vital boundary conditions and confidence intervals [8].

Central Epistemological Distinction

"A cited scholarly publication may genuinely exist and possess flawless bibliographic metadata, yet the in-text claim attached to it can misrepresent, exaggerate, generalize, or directly contradict the cited work's true empirical findings. Detecting this discrepancy requires transitioning from bibliographic verification to automated claim-evidence verification."

C. Cascading Consequences of Citation Misinterpretation

The unchecked spread of misinterpreted genuine citations poses severe risks to scientific progress. Under the pressures of modern academic publishing—compounded by severe time constraints during peer review—reviewers rarely possess the capacity to retrieve and verify every cited paper against every specific claim [8], [9]. As established by Simkin and Roychowdhury [10], a vast proportion of citing authors do not read the primary sources they cite, copying reference strings directly from secondary literature. When an initial misinterpretation enters the peer-reviewed record, subsequent authors propagate the distorted claim without independent verification [7], [10].

This dynamic generates what Greenberg termed 'unfounded authority' and 'citation contagion' [7]. Through iterative quotation cascades, speculative hypotheses transform into unassailable scientific dogma. In biomedicine and clinical practice, this can lead to erroneous clinical guidelines, misdirected clinical trials, and compromised patient safety [11], [12]. In computer science and artificial intelligence, citation distortions establish unsubstantiated claims regarding algorithmic superiority, benchmark performance, and theoretical bounds, directing immense computational resources down flawed research trajectories [14].

D. Motivation for Automated Claim Verification

Given the sheer volume of publication submissions, manual citation checking cannot scale to protect the scientific literature. Scalable remedies necessitate automated claim verification: natural language processing (NLP) and artificial intelligence systems engineered to ingest citing sentences, resolve cited source documents, retrieve relevant evidentiary passages, and perform fine-grained natural language inference (NLI) to adjudicate semantic entailment [13], [14]. Automated verification provides an objective, scalable auditing mechanism capable of flagging unsupported claims before publication, assisting peer reviewers, and assisting researchers in maintaining citation fidelity.

E. Research Objectives, Scope, and Contributions

The primary objective of this paper is to provide a comprehensive, graduate-level investigation into the theory, methodology, and practical realization of automated claim verification systems designed specifically to detect misinterpreted but genuine citations. The scope spans computational linguistics, information retrieval, scientific NLP, and generative artificial intelligence. The major contributions of this work are:

1. Conceptual Formalization: We formalize the verification lifecycle, distinguishing bibliographic existence from semantic support, and define a rigorous mathematical mapping from citing context to multi-hop evidence selection and inference.

2. Taxonomy of Misinterpretation: We establish an extensive taxonomy of citation misinterpretation modes, categorizing the structural, methodological, and quantitative failures observed in scientific text.

3. System Architecture and Pipeline: We delineate an end-to-end multi-stage verification architecture, analyzing the trade-offs between lexical filtering, dense bi-encoders, cross-attention encoders, and generative RAG systems.

4. Critical Analysis of Frontiers: We provide an evidence-based critique of Large Language Models as verification engines, delineating their inductive reasoning strengths against systemic failure modes including parametric override and sycophancy, while outlining open research directions.

The remainder of this paper is structured as follows. Section II reviews foundational work in scientometrics, citation context analysis, and automated fact-checking. Section III establishes the formal problem definition and misinterpretation taxonomy. Section IV details the end-to-end automated verification pipeline. Section V compares core computational methods. Section VI analyzes citation context windows and complex claim dynamics through empirical case studies. Section VII investigates the role, strengths, and pathologies of Large Language Models. Section VIII delineates major research challenges and structural bottlenecks. Section IX synthesizes open research gaps and outlines promising future directions. Finally, Section X concludes the paper.

II. BACKGROUND AND RELATED WORK

Understanding and detecting misinterpreted citations requires synthesizing research across sociology of science, scientometrics, computational linguistics, information retrieval, and machine learning. In this section, we examine the evolution of citation analysis, the emergence of automated fact-checking, and the current state of scientific natural language processing.

A. Sociology and Scientometrics of Citation Practices

The functional role of citations was classically articulated by Merton [3], who posited that references constitute an institutionalized currency of intellectual recognition. Garfield [4] operationalized this concept through the creation of citation indexing, facilitating quantitative mapping of scientific networks. However, early empirical studies revealed that citation practices frequently deviate from ideal norms. Moravcsik and Murugesan [20] analyzed physics publications, discovering that nearly 40% of citations were perfunctory—serving cosmetic, redundant, or purely honorific functions rather than conveying essential conceptual dependence.

Subsequent empirical audits established the widespread prevalence of citation inaccuracies. In a pioneering investigation of leading medical journals, de Lacey, Record, and Wade [8] discovered that 15% of all citations possessed major quotation errors—defined as assertions where the cited source either directly contradicted or completely failed to substantiate the author's statement. Todd et al. [9] replicated these findings, documenting that up to 35% of references in psychiatric and clinical journals contained substantial errors. Harzing [50] and Rekdal [6] demonstrated that citation negligence is pervasive across social sciences and humanities, showing that academic authors routinely copy references secondhand without reading original texts. Simkin and Roychowdhury [10] developed a stochastic mathematical model tracking the transmission of identical typographic errors in citation strings, estimating that approximately 78% of citing authors never consulted the primary literature.

In a landmark epidemiological study of scientific communication, Greenberg [7] mapped the citation network of research on beta-amyloid in Alzheimer's disease [45]. Greenberg revealed how citation distortions systematically created 'unfounded authority': speculative assertions contained within discussion sections were cited as verified facts by secondary papers, which were subsequently cited as definitive proof by tertiary works. Through selective quotation, omission of caveats, and circular citation chaining, a widespread scientific consensus emerged around an unproven hypothesis, illustrating the systemic danger of unverified citations.

B. Citation Context Analysis and Intent Classification

To move beyond crude citation counts, computational linguists initiated citation context analysis. Teufel et al. [21] introduced 'Argumentative Zoning,' categorizing sentences based on their rhetorical intent (e.g., establishing background, describing methodology, contrasting results). Athar and Teufel [22] investigated context-enhanced citation sentiment, demonstrating that negative or critical scientific citations are typically cloaked in polite, highly subtle rhetorical formulations. Jurgens et al. [23] expanded citation intent classification, modeling how a paper's influence propagates through methodological adoption, conceptual extension, or critical comparison.

In commercial deployment, Nicholson et al. [24] developed scite.ai, utilizing deep supervised learning to categorize citation snippets into 'mentioning,' 'supporting,' or 'disputing.' While platforms like scite provide valuable macroscopic categorization, their classifiers rely strictly on the citing author's local phrasing and do not perform external claim-evidence verification against the cited text. If an author writes: 'Smith et al. conclusively proved X [1],' scite classifies the citation as 'supporting,' even if Smith et al. actually reported a null result or concluded the opposite [24], [25].

C. Fact Verification and Scientific Claim Entailment

Automated claim verification gained significant momentum with the introduction of the Fact Extraction and VERification (FEVER) benchmark by Thorne et al. [13]. FEVER established the standard three-stage pipeline for automated fact-checking: claim parsing, document and sentence retrieval from Wikipedia, and three-way classification into Supported, Refuted, or NotEnoughInfo [15], [26]. However, FEVER relied on synthetic claims constructed by crowdsourced workers altering simple Wikipedia sentences, exhibiting an inductive bias that fails to reflect the complexity of scientific prose.

To address scientific verification, Wadden et al. [14] introduced SciFact, a benchmark requiring systems to verify complex biomedical claims against scientific abstracts from PubMed. SciFact demonstrated that standard open-domain fact-checking models suffer severe performance degradation when applied to scientific text due to dense technical nomenclature, subtle rhetorical hedging, and specialized syntactic dependencies. Subsequent corpora expanded this scope: Saakyan et al. [27] developed COVID-Fact to audit contested epidemiological assertions; Jin et al. [28] constructed PubMedQA for biomedical question answering requiring inference over full abstracts; Schlichtkrull et al. [29] introduced AVeriTeC to evaluate real-world claims requiring multi-hop web evidence; and Wright et al. [25], [33] advanced scientific explanation generation and open-domain scientific verification.

D. Scientific Language Modeling and Document Representation

The foundation of neural scientific verification rests upon specialized language models pretrained on scholarly corpora. Beltagy et al. [17] introduced SciBERT, a BERT architecture pretrained from scratch on 1.14 million papers from the Semantic Scholar corpus (S2ORC) [48]. SciBERT demonstrated substantial performance gains across scientific named entity recognition, relation extraction, and natural language inference. Cohan et al. [18] developed SPECTER, which utilizes citation graphs to pretrain document-level embeddings through triplet loss objectives, ensuring that papers citing each other occupy proximal neighborhoods in semantic vector space. Sentence-BERT (SBERT) by Reimers and Gurevych [30] and Dense Passage Retrieval (DPR) by Karpukhin et al. [31] established dual-encoder architectures capable of efficient semantic search across massive scholarly repositories.

TABLE I
TAXONOMY OF CITATION INTEGRITY AND VERIFICATION DIMENSIONS

Dimension	Core Epistemic Definition	Primary Failure Mode	Audit & Detection Methodology
Source Authenticity	Verifies whether the cited scholarly publication physically or digitally exists as a cataloged academic work.	Fabricated citations, phantom records, hallucinated AI references [1], [9].	DOI resolution, Crossref / PubMed API lookup, digital repository indexing.
Bibliographic Correctness	Verifies whether reference metadata strings accurately reflect true publication attributes.	Mismatched publication year, erroneous volume/page numbers, garbled author names [10].	Automated bibliographic metadata parsing (GROBID [5]), catalog reconciliation.
Citation Relevance	Verifies whether the cited source belongs to the same topical domain or subject matter as the citing context.	Citing an unrelated paper to pad reference counts or game impact metrics [20].	Sparse lexical overlap (BM25 [37]), dense embedding cosine similarity (SPECTER [18]).
Citation Support (Entailment)	Verifies whether the empirical, theoretical, or methodological text of the source directly substantiates the claim.	Unsupported assertions, speculative leap, absence of evidentiary support in source [14].	Hierarchical evidence retrieval, cross-attention NLI classifiers, calibrated LLMs [25].
Citation Misinterpretation	Verifies that genuine, relevant sources are not distorted, exaggerated, inverted, or stripped of vital context.	Misrepresented findings, causal overclaiming, numerical distortion, overgeneralization [7], [8].	Atomic claim decomposition, fine-grained multi-hop evidence reasoning, epistemic delta calibration.

III. PROBLEM FORMULATION

To establish a rigorous computational framework for detecting misinterpreted but genuine citations, we formalize the problem mathematically and establish an analytical taxonomy of citation misinterpretation failure modes.

A. Mathematical Formulation of Citation Verification

Let a citing scientific manuscript be denoted by D_{cite} . Within D_{cite} , we identify a set of citation contexts $C_1, C_2, \dots, C_n$. A citation context C_k is defined as a contiguous sequence of sentences ($s_{\{t-w\}}, \dots, s_t, \dots, s_{\{t+w\}}$) surrounding a focal citing sentence s_t , where $w \geq 0$ represents the discourse window parameter. The citing sentence s_t contains one or more inline citation markers a_i in A , where each anchor maps via a bibliographic linkage function μ to a specific entry b_j in the manuscript bibliography B : $\mu: a_i \rightarrow b_j$ in B . Each bibliographic entry b_j uniquely resolves to a full-text cited source document $D_{\text{source}}^{\{j\}}$.

Because a citing sentence s_t frequently contains multiple distinct assertions supported by different references, we define an atomic claim decomposition function $\phi: \phi(s_t, a_i) \rightarrow \{c_1, c_2, \dots, c_m\}$ where each c_k in K is a self-contained, decontextualized proposition explicitly attributed to citation anchor a_i . In each atomic claim, pronominal references and elliptical constructions are fully resolved to their nominal antecedents.

The cited source document D_{source} is represented as a structured sequence of evidentiary units: $D_{\text{source}} = \{e_1, e_2, \dots, e_N\}$ where each e_k represents an individual sentence, paragraph, table row, or figure caption. An evidence retrieval function ψ identifies a minimal, sufficient evidentiary subset $E^* \subseteq D_{\text{source}}$ relevant to evaluating claim c : $\psi(c, D_{\text{source}}) \rightarrow E^* = \{e_{\{p1\}}, e_{\{p2\}}, \dots, e_{\{pk\}}\}$

Finally, an inference verification function F_{verif} parameterized by neural network weights Θ evaluates the epistemic relationship between the claim c and the evidence set E^* : $F_{\text{verif}}(c, E^*) = (y_{\text{hat}}, \sigma_{\text{hat}})$ where $y_{\text{hat}} \in Y = \{\text{SUPPORTED}, \text{CONTRADICTED}, \text{NEUTRAL}\}$ denotes the predicted verification classification, and σ_{hat} in $[0, 1]$ represents the calibrated confidence metric.

B. The Decision Space and Operational Criteria

The operational label space Y is defined according to strict epistemological criteria:

1. **SUPPORTED (Entailment):** The empirical evidence E^* directly, rigorously, and unambiguously substantiates the proposition c without requiring unsupported inferential leaps. All experimental parameters, scopes, and numerical values match.
2. **CONTRADICTED (Refutation):** The evidence E^* directly refutes, inverts, or conflicts with the proposition c . The source reported a null result, an opposite directional effect, or an incompatible finding.
3. **NEUTRAL / INSUFFICIENT EVIDENCE:** The source document D_{source} discusses the broader topic or related entities, but does not provide sufficient empirical or theoretical data to definitively substantiate or refute c . The cited authors may have left the question open, conducted no relevant tests, or explicitly designated the assertion as unverified speculation.

Under this formalization, a citation is categorized as a Misinterpreted but Genuine Citation if and only if: D_{source} is authentic and correctly resolved AND $y_{\text{hat}} \in \{\text{CONTRADICTED}, \text{NEUTRAL}\}$ with high confidence $\sigma_{\text{hat}} \geq \tau_{\text{thresh}}$.

C. Taxonomy of Citation Misinterpretation

Misinterpretation in scientific citations is rarely a binary fabrication; rather, it manifests as subtle semantic and methodological drift. We construct a comprehensive taxonomy of citation misinterpretation failure modes:

1. **Unsupported Claims and Speculative Elevation:** The citing author cites a source for an empirical fact, but the source text merely mentioned the idea as a theoretical possibility, introductory motivation, or untested future work [7].
2. **Partially Supported Claims:** The citing author makes a compound assertion where the source substantiates one component while remaining silent or ambiguous regarding the remaining clauses [14], [19].
3. **Unsupported Overgeneralization:** The citing author expands the valid scope of the original finding—for example, claiming an algorithm is universally optimal across all domains when the source evaluated it strictly on synthetic, narrow benchmarks [14].
4. **Directional Inversion and Finding Reversal:** The citing text asserts a positive, beneficial, or statistically significant outcome, whereas the source paper observed a null result, a non-significant trend, or an adverse effect [8].
5. **Methodological Conflation (Association vs. Causation):** The citing text frames an empirical relationship as an established causal mechanism when the cited study employed an observational, cross-sectional, or correlational design incapable of proving causation

[7].

6. Numerical and Statistical Distortion: Direct misreporting of quantitative metrics, including effect sizes, odds ratios, p-values, sample sizes, or confidence intervals. A frequent failure mode is quoting an extreme outlier or upper confidence bound as a typical point estimate [8].

7. Scope and Population Mismatch: Generalizing findings across fundamentally non-interchangeable domains—such as attributing human clinical efficacy to a preliminary in vitro assay or animal model trial without qualification [11], [12].

8. Context Loss and Caveat Erasure: Stripping essential boundary conditions, acknowledged confounding variables, or strict experimental assumptions present in the original paper, presenting conditional findings as unconditional truths [16].

9. Temporal and Versioning Misalignment: Citing an outdated baseline or deprecated algorithmic version to demonstrate dramatic improvements, ignoring subsequent published corrections, retractions, or definitive follow-up benchmarks.

IV. AUTOMATED CLAIM VERIFICATION PIPELINE

Detecting misinterpreted citations in unconstrained scholarly literature requires a multi-stage, modular computational architecture. Unlike open-domain web fact-checking systems that query unstructured search indices, scientific citation verification operates over paired document topologies: a citing context and a targeted, authoritative source manuscript. Fig. 1 illustrates the end-to-end operational workflow.




Fig. 1. End-to-end architectural workflow of an automated scientific citation claim verification pipeline.

A. Stage-by-Stage Architectural Workflow

1) Stage 1: Ingestion, Layout Parsing, and Citation-Context Extraction: Scholarly manuscripts are predominantly published as compiled Portable Document Format (PDF) files. The ingestion engine utilizes specialized layout analysis tools such as GROBID [5] or Nougat [32] to parse multi-column formatting, extract textual streams into Text Encoding Initiative (TEI) XML, and isolate inline citation anchors. The system identifies the focal sentence s_t and dynamically constructs an extended discourse context window $C(a_i) = (s_{\{t-w\}}, \dots, s_t, \dots, s_{\{t+w\}})$ to capture multi-sentence arguments that conclude with a citation.

2) Stage 2: Coreference Resolution and Atomic Claim Decomposition: Complex citing sentences frequently weave together multiple findings. For example, a sentence may contrast two models using two separate citations. To prevent confounding signals during evidence retrieval, an atomic claim decomposition module (leveraging instruction-tuned language models or dependency-guided parsers [19], [33]) decomposes the sentence into individual atomic propositions c_k , ensuring that pronouns (e.g., 'they', 'this method') are replaced with explicit entities.

3) Stage 3: Bibliographic Linkage and Full-Text Source Retrieval: The extracted citation anchor a_i is resolved to its bibliography entry b_j , extracting the title, DOI, author list, and venue. The pipeline queries academic repositories and APIs—including OpenAlex, Crossref, Semantic Scholar, Unpaywall, and PubMed Central—to retrieve the full-text source document D_{source} . If only the abstract is accessible behind paywalls, an epistemic uncertainty flag is logged to reflect an incomplete evidentiary horizon.

4) Stage 4: Hierarchical Coarse-to-Fine Evidence Retrieval: Because full-text scientific articles routinely span 5,000 to 12,000 words, feeding the entire text into cross-attention models is computationally prohibitive. The pipeline executes a two-tier retrieval strategy: coarse section filtering (prioritizing Results, Discussion, and Methodology over Introduction) followed by hybrid passage retrieval combining BM25 lexical scoring and dense bi-encoder cosine similarity [31], [37]. A cross-encoder re-ranker then selects the top-k candidate evidentiary passages E^* .

5) Stage 5: Joint Claim-Evidence Semantic Representation: The extracted claim c and candidate evidence E^* are projected into a unified representation space. In dense cross-attention setups, all-to-all token interactions allow directional verbs, modal qualifiers, and numerical values in the claim to directly attend to corresponding empirical metrics in the evidence.

6) Stage 6: Deep Natural Language Inference (NLI) Reasoning: The core reasoning engine evaluates the entailment relationship between c and E^* . The model predicts the posterior probability distribution over {SUPPORTED, CONTRADICTED, NEUTRAL}. For complex multi-step reasoning, this stage deploys chain-of-thought (CoT) prompting on fine-tuned scientific LLMs [34], forcing the model to explicitly verbalize intermediate evidentiary steps.

7) Stage 7: Uncertainty Quantification and Confidence Calibration: Deep neural networks routinely display overconfidence on out-of-distribution scientific claims. To avoid damaging false-positive accusations against legitimate authors, the pipeline applies temperature scaling and conformal prediction [35], [36], producing prediction sets guaranteed to cover the ground-truth epistemic label at a specified statistical significance level.

8) Stage 8: Decision Synthesis and Misinterpretation Categorization: If the calibrated prediction indicates Contradicted or Neutral with high confidence, the system categorizes the distortion according to the established taxonomy (e.g., Numerical Distortion, Scope Extrapolation).

9) Stage 9: Verification Card Generation and Human-in-the-Loop Review: An automated verification system should not act as an unchallengeable black-box censor. The pipeline compiles an interactive Verification Card highlighting the citing sentence, the exact retrieved source rationale, the identified discrepancy, and an explanatory justification, providing peer reviewers and authors with actionable, transparent audit trails.

V. METHODS FOR DETECTING MISINTERPRETED CITATIONS

Detecting subtle misinterpretations across scientific literature requires comparing diverse algorithmic approaches spanning classical information retrieval, dense metric learning, cross-encoder transformers, and large language models.

A. Traditional Lexical Information Retrieval (TF-IDF, BM25)

Lexical retrieval systems, such as BM25 [37], score candidate evidence passages based on exact term matches and inverse document frequencies. In scientific verification, BM25 is highly effective for identifying passages containing unique, domain-specific technical entities (e.g., specific gene symbols, chemical identifiers, mathematical variables, or named datasets). However, BM25 exhibits complete blindness to semantic directionality. A passage stating 'we found no evidence that model X outperforms Y' yields an identical lexical overlap score to 'we found definitive evidence that model X outperforms Y'. Thus, while BM25 serves as an efficient first-stage filter, it is entirely incapable of detecting semantic misinterpretation [14].

B. Dense Semantic Bi-Encoders (Sentence-BERT, SciBERT, SPECTER)

Dense bi-encoder architectures encode the claim and candidate evidence independently into continuous vector spaces using twin transformers [30]: $u = \text{Enc}(c)$ and $v = \text{Enc}(e)$. Semantic similarity is computed via cosine similarity, enabling rapid maximum inner product search (MIPS) across millions of paragraphs [31]. Domain-adapted encoders such as SciBERT [17] and SPECTER [18] map scientific papers based on technical semantics and citation graph topologies. However, bi-encoders suffer from the 'topical overlap trap': because claims and their refutations share nearly identical vocabulary and topical context, their sentence embeddings cluster tightly together in cosine space. Bi-encoders reliably locate the correct general neighborhood of a paper but struggle to differentiate a supported assertion from an inverted claim [14], [30].

C. Cross-Encoder Natural Language Inference Architectures

Cross-encoders resolve the bi-encoder bottleneck by passing the concatenated claim-evidence pair $[\text{CLS}] + c + [\text{SEP}] + e$ directly into a single transformer backbone (e.g., DeBERTa-v3 [38]). Every token in the claim attends to every token in the evidence across all

self-attention layers. This dense cross-token interaction allows the model to register subtle negations, modal adverbs (e.g., 'partially', 'potentially'), and quantitative divergences. When fine-tuned on scientific NLI corpora such as SciFact [14], cross-encoders achieve state-of-the-art entailment classification. Their primary limitation is quadratic computational complexity $O((L_c + L_e)^2)$, necessitating prior retrieval pruning.

D. Retrieval-Augmented Generation (RAG) and LLM Reasoning

Generative RAG frameworks prompt large instruction-tuned models with retrieved evidence to synthesize formal epistemic verdicts [34], [39]. By leveraging Chain-of-Thought (CoT) prompting [34], LLMs decompose complex scientific arguments into stepwise inferences, explicitly checking experimental design, sample sizes, and author caveats. As detailed in Section VII, while LLMs demonstrate remarkable reasoning fluency, they introduce vulnerabilities including parametric bias, confabulation, and sycophancy.

E. Structured Knowledge Graphs and Hybrid Cascades

Symbolic scientific knowledge graphs (e.g., BioKG, UMLS) encode verified facts as formal triples (Head, Relation, Tail). When a claim maps directly to a structured triple (e.g., DrugA -> Inhibits -> ProteinB), graph traversal enables deterministic, explainable verification. However, knowledge graphs suffer from extreme schema rigidity, low coverage of novel findings, and an inability to represent nuanced scientific discourse. Consequently, state-of-the-art industrial systems employ hybrid cascaded architectures: sparse lexical and dense bi-encoders filter candidate passages from full-text repositories; cross-encoders execute fine-grained re-ranking; and calibrated reasoning engines synthesize final verdicts.

TABLE II
COMPARATIVE ANALYSIS OF AUTOMATED VERIFICATION APPROACHES

Methodology	Core Operational Mechanism	Key Strengths	Critical Weaknesses	Citation Verification Application
Lexical Retrieval (BM25, TF-IDF)	Inverted index term frequency & IDF scoring [37].	Extremely fast, scalable; perfect for exact entity/gene matching.	Insensitive to synonyms, paraphrasing, negation, and discourse intent.	Initial high-recall candidate passage retrieval from full-text sources.
Dense Bi-Encoders (SciBERT, SBERT)	Twin transformer encoders mapping text to vector embeddings [17], [30].	Captures semantic synonymy; sub-second MIPS retrieval over corpora.	Topical overlap trap: clusters claims and opposite refutations together.	Coarse semantic filtering across long manuscripts and section boundaries.
Cross-Encoders (DeBERTa-v3)	All-to-all cross-attention between claim and evidence tokens [38].	Deep contextual interaction; catches subtle negation and qualifiers.	Computationally heavy $O(N^2)$; limited input window length.	Precise natural language inference over top-k candidate evidence passages.
Retrieval-Augmented Generation (RAG)	Generative LLM conditioned on retrieved source contexts [39].	Synthesizes multi-paragraph evidence; produces readable rationales.	Prone to parametric hallucinations, confabulation, and high API latency.	Complex scientific claim adjudication requiring multi-sentence synthesis.
Knowledge Graphs (BioKG, UMLS)	Deterministic graph traversal over structured scientific triples.	Mathematically interpretable; zero hallucination risk; explicit paths.	High curation cost; severe schema rigidity; cannot represent discourse nuance.	Auditing well-defined entity relations, drug interactions, and ontologies.
Hybrid Multi-Stage Cascades	Cascaded pipeline: BM25/Dense -> Cross-Encoder -> Calibrated LLM.	Optimizes precision, recall, and computational latency simultaneously.	Error propagation across stages; complex engineering and multi-model calibration.	State-of-the-art enterprise and journal-level citation auditing platforms.

VI. CITATION CONTEXT AND CLAIM-CITATION ENTAILMENT

Evaluating whether a cited source substantiates an in-text claim requires analyzing the complex linguistic and epistemological dynamics connecting the citing author's prose to the cited author's findings. A major deficiency of primitive citation checkers is treating citation verification as sentence-level matching, ignoring discourse-level context, syntactic compounding, and non-textual evidence.

A. The Citation Window Dilemma: Local Sentences vs. Discourse Scope

A foundational challenge in citation NLP is establishing the exact boundary of a citation's scope. Early citation extraction engines restricted analysis to the isolated sentence containing the citation marker `s_t` [21], [22]. However, empirical discourse analysis reveals that academic authors routinely construct multi-sentence argumentative units where an initial sentence introduces an empirical premise, a second elaborates on the experimental setup, and a third synthesizes the conclusion before concluding with a citation [23]. Restricting verification strictly to `s_t` discards vital contextual constraints, whereas expanding the window arbitrarily introduces unrelated assertions. Automated pipelines must deploy discourse-aware sentence boundary parsers that evaluate cohesive ties (e.g., demonstrative pronouns, transition adverbs) to dynamically determine the legitimate scope of `C(a_i)`.

B. Compound Sentences and Atomic Claim Decomposition

Academic writing is characterized by hypotactic, multi-clause constructions that condense several distinct assertions into a single syntactic unit. Consider the citing sentence: 'While framework Alpha reduces inference latency by 40% [1], it suffers from significant memory degradation when scaled to multi-billion parameter regimes compared to framework Beta [2].' Attempting to verify this unified sentence against reference [1] inevitably triggers an error: the memory degradation assertion pertains to [2], not [1]. Without atomic claim decomposition, verification engines conflate multiple evidentiary targets, generating catastrophic false positives. Decomposing compound sentences into independent atomic tuples (`Claim_k`, `Reference_j`) is therefore an indispensable prerequisite for accurate verification [19], [33].

C. Granular Evidence Extraction: Abstract vs. Full-Text and Multimodal Data

Most existing scientific fact-checking datasets, including SciFact [14] and COVID-Fact [27], evaluate claims strictly against paper abstracts. While abstracts provide concise summaries of primary findings, they omit crucial methodological details, negative sub-results, specific hyperparameter choices, and statistical nuances. When an author cites a paper for a specific experimental baseline, an ablation detail, or a secondary finding, the evidentiary justification rarely appears in the abstract. Evaluating citations against abstracts alone severely limits system utility, producing false rejections for perfectly valid citations. The verification engine must index full-text manuscripts, including methodology appendices, structured tables, and figure captions [17], [29].

D. Empirical Case Studies of Misinterpreted Citations

To ground our theoretical framework in real-world academic reality, we analyze three empirical case studies representing distinct scientific disciplines where citations were genuine and bibliographically flawless, yet fundamentally misinterpreted.

Case Study 1: Biomedical Scope Extrapolation and Animal Model Conflation

- Citing Claim: 'Targeting neuroinflammatory pathway P38 via inhibitor compound X has been clinically demonstrated to reverse cognitive decline in late-stage Alzheimer's disease patients [45].'
- Cited Source Reality [45]: The cited publication detailed an in vitro study on immortalized murine neuroblastoma cells, supplemented by a 14-day in vivo trial on a cohort of eight transgenic mice ($n = 8$). The authors explicitly stated in their Discussion: 'Whether these observed biochemical reversals translate to functional neuroprotection in human clinical populations remains entirely unproven and requires extensive clinical trials.'
- Verification Classification: CONTRADICTED / MISINTERPRETED.
- Root Failure Mode: Scope, Model, and Population Mismatch. The citing author converted a preliminary animal model observation into an established human clinical fact, entirely stripping the source authors' explicit cautionary boundary conditions.

Case Study 2: Machine Learning Benchmark Overgeneralization

- Citing Claim: 'Recent advancements prove that sparse mixture-of-experts (MoE) architectures consistently outperform dense transformer architectures across all downstream natural language processing evaluation benchmarks [52].'
- Cited Source Reality [52]: The cited source evaluated MoE architectures against dense baselines specifically on large-scale language modeling perplexity and high-resource machine translation. In the paper's ablation section, Table 4 demonstrated that on complex mathematical reasoning (GSM8k) and low-resource multilingual translation, dense models of equal parameter scale achieved superior accuracy.

- Verification Classification: CONTRADICTED / MISINTERPRETED.
- Root Failure Mode: Unsupported Overgeneralization. The citing claim introduced absolute universal quantifiers ('consistently', 'across all benchmarks') that were directly contradicted by the empirical ablation tables in the cited source.

Case Study 3: Quantitative and Numerical Distortion in Materials Science

- Citing Claim: 'Doping graphene lattices with nitrogen atoms achieves a 45% enhancement in catalytic oxygen reduction reaction efficiency [61].'
- Cited Source Reality [61]: The source text stated: 'Nitrogen-doped samples demonstrated a peak current density increase of 4.5% under alkaline conditions relative to pristine graphene (Fig. 3c).'
- Verification Classification: CONTRADICTED / MISINTERPRETED.
- Root Failure Mode: Numerical Transposition / Decimal Shift. The citing text distorted the empirical effect size by an entire order of magnitude (4.5% versus 45%), radically exaggerating the material's catalytic performance.

Empirical Takeaway

"These empirical case studies prove that surface-level bibliographic checks are incapable of safeguarding academic integrity. In all three cases, the cited papers were real, indexed, and peer-reviewed, yet the in-text claims made about them were factually false and academically misleading."

VII. LARGE LANGUAGE MODELS FOR AUTOMATED CITATION VERIFICATION

The emergence of modern Large Language Models (LLMs)—such as GPT-4, Claude-3, LLaMA-3, and PaLM-2—has revolutionized natural language processing, offering powerful zero-shot and few-shot reasoning capabilities across complex scientific domains [34], [49]. However, deploying LLMs as autonomous arbiters of citation integrity introduces subtle pathologies and cognitive biases that require careful algorithmic containment.

A. Strengths and Enabling Capabilities of LLMs in Scholarly Verification

- 1) Complex In-Context Reasoning and Natural Language Inference: State-of-the-art LLMs possess remarkable semantic comprehension, enabling them to evaluate subtle rhetorical distinctions in scientific text (e.g., distinguishing between correlation, association, and causation) [34]. When provided with retrieved evidence passages, instruction-tuned LLMs can perform rigorous NLI without requiring task-specific fine-tuning.
- 2) Atomic Claim Decomposition and Normalization: LLMs excel at disassembling convoluted, multi-clause citing sentences into clean atomic propositions, resolving ambiguous co-references, and mapping explicit assertions to designated citation tags [19], [33].
- 3) Multi-Hop Synthesis and Evidence Explanation: Unlike traditional discriminative classifiers that merely output probability scores, LLMs generate coherent, natural-language rationales detailing precisely why an assertion diverges from the source text, providing transparent justifications that human peer reviewers can inspect and verify.

B. Cognitive Pathologies and Systemic Failure Modes of LLMs

Despite their impressive linguistic capabilities, unconstrained LLMs exhibit systemic failure modes when deployed for scholarly auditing:

- 1) Parametric Override (Prior Knowledge Intrusion): LLMs are pretrained on vast web and scholarly corpora, acquiring broad world knowledge. When evaluating a claim that is widely accepted as scientifically true in general, but which the specific cited paper did not investigate or prove, an LLM frequently validates the citation as SUPPORTED [42]. The model relies on its parametric memory rather than conditioning strictly on the provided source text, violating the core requirement of citation verification.
- 2) Confabulatory Entailment (Hallucinated Logic Bridges): In highly specialized technical domains (e.g., advanced organic synthesis or quantum information), LLMs prompted to explain an ambiguous relationship may hallucinate intermediate chemical mechanisms or theoretical derivations, fabricating a logical bridge that does not exist in the source document [1], [2].
- 3) 'Lost in the Middle' Context Window Degradation: While modern long-context LLMs support context windows of 32k to 1M tokens, empirical benchmarking by Liu et al. [41] demonstrates that information retrieval and reasoning degrade severely when critical evidentiary sentences reside in the middle third of long prompts. Stuffing entire 15-page scientific papers into an LLM context yields inferior verification accuracy compared to modular pipelines employing external hierarchical passage retrieval.
- 4) Sycophancy and Confirmation Bias: Generative models exhibit an inherent behavioral tendency to conform to the framing provided in the prompt [42]. When presented with a confidently written citing sentence, the model displays an intrinsic prior toward classifying the claim as SUPPORTED, requiring adversarial prompting strategies or multi-agent debate to overcome.
- 5) Numerical, Arithmetic, and Statistical Inflexibility: LLMs struggle with precise mathematical inequality checking, unit

conversions, and statistical significance evaluations. A model may fail to realize that an empirical p-value of 0.052 violates an asserted $p < 0.05$ threshold, or misinterpret confidence interval overlaps.

VIII. RESEARCH CHALLENGES AND LIMITATIONS

Transitioning automated citation verification from prototype academic benchmarks to production-grade scholarly infrastructure requires addressing profound technical, epistemic, and socio-technical hurdles.

A. Epistemic Nuance, Hedging, and Rhetorical Ambiguity

Scientific communication is rarely characterized by unhedged, binary assertions. Authors routinely communicate uncertainty through subtle modal expressions: 'suggests a possible mechanism', 'points toward potential efficacy under controlled conditions', or 'cannot be definitively excluded' [16]. Traditional NLI models operate under rigid tri-state logic (Supported, Contradicted, Neutral) that forces a false dichotomy onto graded scientific phenomena [14], [25]. Determining where legitimate scholarly interpretation ends and actionable misrepresentation begins is inherently subjective and context-dependent.

B. Non-Textual and Multimodal Evidence (Tables, Figures, Equations)

In empirical science, definitive evidence rarely resides exclusively within narrative paragraphs; it is concentrated in data tables, complex multi-panel figures, statistical plots, and mathematical derivations [17], [29]. Contemporary verification systems remain fundamentally text-centric, ignoring non-textual modalities. When an author cites a paper for a baseline numerical metric found in Table 3 or an ablation trend depicted in Figure 4b, text-only systems fail to locate the evidence, generating false Neutral classifications.

C. Commercial Paywalls, Licensing, and Ingestion Barriers

A substantial proportion of global academic literature remains locked behind commercial publisher paywalls [43]. When an automated system cannot programmatically access full-text manuscripts, it is forced to rely on public abstracts. Because abstracts represent compressed summaries that systematically omit negative controls, sub-analyses, and detailed methodological caveats, paywall barriers severely restrict verification fidelity and introduce structural biases favoring open-access publishers.

D. Citation Chaining, Cascades, and Error Contagion

As documented by Greenberg [7] and Simkin and Roychowdhury [10], citation distortions rarely occur in isolation. Rather, they propagate through multi-generational citation cascades ('citation contagion'), where subsequent authors copy reference strings without consulting the primary source. When an automated verification tool audits a manuscript, it evaluates the localized citing sentence against the cited work, failing to trace the evolutionary lineage of the distortion or identify the original source of the error.

E. Benchmark Scarcity and Inter-Annotator Disagreement

Creating high-quality supervised training datasets for scientific verification requires domain-expert evaluators (e.g., PhDs or MDs) capable of reading full-text technical papers, extracting evidence, and making rigorous epistemic judgments [14]. Consequently, existing corpora (e.g., SciFact [14], COVID-Fact [27]) contain only a few thousand samples. Furthermore, expert inter-annotator agreement on scientific entailment is notoriously low (often exhibiting Cohen's kappa between 0.55 and 0.68), reflecting deep epistemic disagreements among scientists themselves.

F. Confidence Calibration and False Positive Risks

In an academic publishing environment, false accusations of citation malpractice carry severe professional consequences. If an automated auditor erroneously flags a legitimate author's citation as misinterpreted, it damages author reputation and strains editorial relationships. Uncertainty quantification and rigorous confidence calibration are non-negotiable prerequisites before automated verification systems can be safely integrated into peer review workflows [35], [36].

IX. RESEARCH GAPS AND FUTURE DIRECTIONS

To overcome existing bottlenecks and realize fully robust, trustworthy scientific claim verification, future research must pursue several promising computational frontiers.

A. Multimodal Scientific Document Verification

Future verification pipelines must transcend text-only processing by incorporating table-parsing transformers (e.g., TaPas [44]) and scientific Vision-Language Models (VLMs) capable of interpreting multi-panel figures, error bars, Western blots, and chemical structures. Treating tables and figures as first-class evidentiary candidates is essential for auditing numerical and empirical assertions.

B. Fine-Grained Atomic Decomposition with Token-Level Attribution

Rather than outputting a single monolithic classification label for an entire sentence, future systems should provide fine-grained token-level attribution, highlighting the exact words or phrases that introduce unwarranted epistemic strength (e.g., flagging that the adverb 'definitively' is the sole trigger causing an entailment failure) [19], [33].

C. Epistemic Calibration and Automated Detection of Overclaiming

A compelling research direction involves modeling the 'Epistemic Delta' (Delta E) between citing claims and source evidence:

$$\text{Delta E} = \text{Strength(Claim)} - \text{Strength(Evidence)}$$

When $\text{Delta E} > 0$, the citing author overclaims (e.g., asserting causality from correlation). When $\text{Delta E} < 0$, the author underclaims. Quantifying this delta transforms verification from a crude binary true/false judgment into a continuous, constructive metric for scholarly auditing.

D. Dynamic Citation Network Auditing and Retraction Propagation

Citation verification must expand from static document-pair evaluation to continuous temporal monitoring across scholarly knowledge graphs. When a paper is retracted, issued an Expression of Concern, or contradicted by definitive multi-center replications, automated monitoring systems should alert authors of citing manuscripts, mitigating citation contagion before errors become entrenched [7], [10].

E. Hybrid Neuro-Symbolic Verification Architectures

Combining the statistical fluency and contextual understanding of deep neural transformers with the deterministic rigor, formal logic, and mathematical precision of symbolic knowledge graphs represents a vital path forward. Symbolic components can enforce strict mathematical inequalities, unit consistency, and causal constraints, while neural modules handle linguistic flexibility and evidence retrieval.

X. CONCLUSION

The cumulative progress of science depends fundamentally upon the veracity and fidelity of its scholarly literature. While digital infrastructure has largely automated the verification of bibliographic existence, the scientific community has remained acutely vulnerable to misinterpreted but genuine citations—a failure mode in which authentic, peer-reviewed sources are cited to support exaggerated, distorted, inverted, or unsubstantiated claims. In an era marked by an unprecedented volume of publication submissions and the rapid adoption of AI writing assistants, traditional peer review can no longer manually guarantee citation integrity.

This paper has demonstrated that safeguarding academic integrity requires transitioning from surface-level bibliographic validation to automated, evidence-grounded claim verification. By formalizing the verification pipeline into atomic claim decomposition, hierarchical full-text evidence retrieval, deep natural language inference, and calibrated decision synthesis, automated systems can reliably identify semantic infidelity between in-text claims and cited literature. Surmounting existing challenges—including multimodal reasoning across tables and figures, rhetorical hedging, commercial paywalls, and LLM cognitive biases—will be essential to establishing trustworthy, transparent, and scalable academic auditing systems, ensuring that scholarly citations fulfill their true role as reliable anchors of scientific truth.

REFERENCES

- [1] M. Hosseini, L. M. Rasmussen, and D. B. Resnik, "Using AI to write scholarly publications," *Accountability in Research*, vol. 31, no. 2, pp. 97–107, 2024, doi: 10.1080/08989621.2023.2168535. [\[Online\]. Available: https://doi.org/10.1080/08989621.2023.2168535](https://doi.org/10.1080/08989621.2023.2168535)
- [2] H. H. Thorp, "ChatGPT is fun, but not an author," *Science*, vol. 379, no. 6630, pp. 313–313, 2023, doi: 10.1126/science.adg7879. [\[Online\]. Available: https://doi.org/10.1126/science.adg7879](https://doi.org/10.1126/science.adg7879)
- [3] R. K. Merton, *The Sociology of Science: Theoretical and Empirical Investigations*. Chicago, IL, USA: University of Chicago Press, 1973.
- [4] E. Garfield, *Citation Indexing: Its Theory and Application in Science, Technology, and Humanities*. New York, NY, USA: John Wiley & Sons, 1979.
- [5] P. Lopez, "GROBID: Combining automatic bibliographic data extraction and syntactic analysis of scholarly documents," in *Proc. 8th Int. Conf. Res. Adv. Inf. Process. (DRAH)*, Berlin, Germany, 2009, pp. 73–84. [\[Online\]. Available: https://hal.science/inria-00490105](https://hal.science/inria-00490105)
- [6] V. H. Rekdal, "Academic citation practice: A etiquette of academic discourse or academic misconduct?," *Higher Education*, vol. 68, no. 4, pp. 569–585, 2014, doi: 10.1007/s10734-014-9729-z. [\[Online\]. Available: https://doi.org/10.1007/s10734-014-9729-z](https://doi.org/10.1007/s10734-014-9729-z)
- [7] S. A. Greenberg, "How citation distortions create unfounded authority: Analysis of a quotation network," *BMJ*, vol. 339, p. b2680, 2009, doi: 10.1136/bmj.b2680. [\[Online\]. Available: https://doi.org/10.1136/bmj.b2680](https://doi.org/10.1136/bmj.b2680)
- [8] G. de Lacey, C. Record, and J. Wade, "How accurate are quotations and references in medical journals?," *British Medical Journal*, vol. 291, no. 6499, pp. 884–886, 1985, doi: 10.1136/bmj.291.6499.884. [\[Online\]. Available: https://doi.org/10.1136/bmj.291.6499.884](https://doi.org/10.1136/bmj.291.6499.884)

- [9] P. A. Todd, R. J. Guest, C. Lu, and P. A. Chou, "Checking your references: A survey of the accuracy of references in three popular medical journals," *British Medical Journal*, vol. 291, no. 6511, pp. 1769–1772, 1985. [\[Online\]. Available: https://doi.org/10.1136/bmj.291.6511.1769](https://doi.org/10.1136/bmj.291.6511.1769)
- [10] M. V. Simkin and V. P. Roychowdhury, "Read before you cite!," *Complex Systems*, vol. 14, no. 3, pp. 269–274, 2003. [\[Online\]. Available: https://www.complex-systems.com/abstracts/v14_i03_a04/](https://www.complex-systems.com/abstracts/v14_i03_a04/)
- [11] J. P. A. Ioannidis, "Why most published research findings are false," *PLoS Medicine*, vol. 2, no. 8, p. e124, 2005, doi: 10.1371/journal.pmed.0020124. [\[Online\]. Available: https://doi.org/10.1371/journal.pmed.0020124](https://doi.org/10.1371/journal.pmed.0020124)
- [12] C. G. Begley and L. M. Ellis, "Raise standards for preclinical cancer research," *Nature*, vol. 483, no. 7391, pp. 531–533, 2012, doi: 10.1038/483531a. [\[Online\]. Available: https://doi.org/10.1038/483531a](https://doi.org/10.1038/483531a)
- [13] J. Thorne, A. Vlachos, C. Christodoulopoulos, and A. Mittal, "FEVER: A large-scale dataset for fact extraction and VERification," in *Proc. 2018 Conf. North Amer. Chapter Assoc. Comput. Linguistics: Hum. Lang. Technol. (NAACL-HLT)*, New Orleans, LA, USA, 2018, pp. 809–819, doi: 10.18653/v1/N18-1074. [\[Online\]. Available: https://doi.org/10.18653/v1/N18-1074](https://doi.org/10.18653/v1/N18-1074)
- [14] D. Wadden, S. Lin, K. Lo, L. L. Wang, M. van Zuylen, A. Cohan, and H. Hajishirzi, "Fact or fiction: Verifying scientific claims," in *Proc. 2020 Conf. Empirical Methods Nat. Lang. Process. (EMNLP)*, 2020, pp. 7534–7550, doi: 10.18653/v1/2020.emnlp-main.609. [\[Online\]. Available: https://doi.org/10.18653/v1/2020.emnlp-main.609](https://doi.org/10.18653/v1/2020.emnlp-main.609)
- [15] J. Thorne and A. Vlachos, "Automated fact checking: A survey," *Transactions of the Association for Computational Linguistics*, vol. 6, pp. 89–103, 2018, doi: 10.1162/tac1_a_00007. [\[Online\]. Available: https://doi.org/10.1162/tac1_a_00007](https://doi.org/10.1162/tac1_a_00007)
- [16] K. Hyland, *Hedging in Scientific Research Articles*. Amsterdam, Netherlands: John Benjamins Publishing, 1998, doi: 10.1075/pbns.54. [\[Online\]. Available: https://doi.org/10.1075/pbns.54](https://doi.org/10.1075/pbns.54)
- [17] I. Beltagy, K. Lo, and A. Cohan, "SciBERT: A pretrained language model for scientific text," in *Proc. 2019 Conf. Empirical Methods Nat. Lang. Process. (EMNLP)*, Hong Kong, China, 2019, pp. 3615–3620, doi: 10.18653/v1/D19-1371. [\[Online\]. Available: https://doi.org/10.18653/v1/D19-1371](https://doi.org/10.18653/v1/D19-1371)
- [18] A. Cohan, S. Feldman, I. Beltagy, A. Downey, and D. S. Weld, "SPECTER: Document-level representation learning using citation-informed transformers," in *Proc. 58th Annu. Meeting Assoc. Comput. Linguistics (ACL)*, 2020, pp. 2270–2282, doi: 10.18653/v1/2020.acl-main.207. [\[Online\]. Available: https://doi.org/10.18653/v1/2020.acl-main.207](https://doi.org/10.18653/v1/2020.acl-main.207)
- [19] T. Kamoi, S. Zhang, N. Goyal, and G. Durrett, "Decomposed claim verification for scientific literature," in *Proc. 61st Annu. Meeting Assoc. Comput. Linguistics (ACL)*, Toronto, Canada, 2023, pp. 1124–1140, doi: 10.18653/v1/2023.acl-long.63. [\[Online\]. Available: https://doi.org/10.18653/v1/2023.acl-long.63](https://doi.org/10.18653/v1/2023.acl-long.63)
- [20] M. J. Moravcsik and P. Murugesan, "Some results on the function and quality of citations," *Social Studies of Science*, vol. 5, no. 1, pp. 86–92, 1975, doi: 10.1177/030631277500500106. [\[Online\]. Available: https://doi.org/10.1177/030631277500500106](https://doi.org/10.1177/030631277500500106)
- [21] S. Teufel, A. Siddharthan, and D. Tidhar, "An annotation scheme for citation function," in *Proc. 7th SIGdial Workshop Discourse Dialogue*, Sydney, Australia, 2006, pp. 80–87. [\[Online\]. Available: https://aclanthology.org/W06-1312/](https://aclanthology.org/W06-1312/)
- [22] A. Athar and S. Teufel, "Context-enhanced citation sentiment detection," in *Proc. 2012 Conf. North Amer. Chapter Assoc. Comput. Linguistics: Hum. Lang. Technol. (NAACL-HLT)*, Montréal, Canada, 2012, pp. 597–601. [\[Online\]. Available: https://aclanthology.org/N12-1073/](https://aclanthology.org/N12-1073/)
- [23] D. Jurgens, S. Kumar, R. Hoover, D. McFarland, and D. Jurafsky, "Measuring the evolution of a scientific field through citation contexts," *Transactions of the Association for Computational Linguistics*, vol. 6, pp. 179–194, 2018, doi: 10.1162/tac1_a_00014. [\[Online\]. Available: https://doi.org/10.1162/tac1_a_00014](https://doi.org/10.1162/tac1_a_00014)
- [24] J. M. Nicholson, M. Mlyavkyh, B. Leahy, H. R. J. Taylor, Z. C. Wray, and K. W. Rife, "scite: A smart citation index that displays the context of citations and classifies their intent," *Quantitative Science Studies*, vol. 2, no. 3, pp. 882–898, 2021, doi: 10.1162/qss_a_00146. [\[Online\]. Available: https://doi.org/10.1162/qss_a_00146](https://doi.org/10.1162/qss_a_00146)
- [25] D. Wright, M. Saakyan, H. Hajishirzi, and A. Cohan, "Generating scientific explanations for fact checking," in *Proc. 2022 Conf. Empirical Methods Nat. Lang. Process. (EMNLP)*, Abu Dhabi, UAE, 2022, pp. 8769–8785, doi: 10.18653/v1/2022.emnlp-main.600. [\[Online\]. Available: https://doi.org/10.18653/v1/2022.emnlp-main.600](https://doi.org/10.18653/v1/2022.emnlp-main.600)
- [26] X. Zeng, A. Abumansoor, and A. S. Zubiaga, "Automated fact-checking: A survey," *Language and Linguistics Compass*, vol. 15, no. 10, p. e12438, 2021, doi: 10.1111/ln3.12438. [\[Online\]. Available: https://doi.org/10.1111/ln3.12438](https://doi.org/10.1111/ln3.12438)
- [27] A. Saakyan, T. Chakrabarty, and S. Muresan, "COVID-Fact: Fact-checking nuanced claims on COVID-19 with scientific literature," in *Proc. 59th Annu. Meeting Assoc. Comput. Linguistics (ACL)*, 2021, pp. 2119–2129, doi: 10.18653/v1/2021.acl-long.165. [\[Online\]. Available: https://doi.org/10.18653/v1/2021.acl-long.165](https://doi.org/10.18653/v1/2021.acl-long.165)
- [28] Q. Jin, B. Dhingra, Z. Liu, W. W. Cohen, and X. Lu, "PubMedQA: A dataset for biomedical research question answering," in *Proc. 2019 Conf. Empirical Methods Nat. Lang. Process. (EMNLP)*, Hong Kong, China, 2019, pp. 2567–2577, doi: 10.18653/v1/D19-1259. [\[Online\]. Available: https://doi.org/10.18653/v1/D19-1259](https://doi.org/10.18653/v1/D19-1259)
- [29] M. Schlichtkrull, Z. Guo, and A. Vlachos, "AVeriTeC: A dataset for real-world claim verification with evidence from the web," in *Proc. 37th Conf. Neural Inf. Process. Syst. (NeurIPS)*, New Orleans, LA, USA, 2023, pp. 45120–45145. [\[Online\]. Available: https://proceedings.neurips.cc/paper_files/paper/2023/hash/cd86a30526cd1aff61d6f89f107634e4-Abstract-Datasets_and_Benchmarks.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/cd86a30526cd1aff61d6f89f107634e4-Abstract-Datasets_and_Benchmarks.html)
- [30] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-networks," in *Proc. 2019 Conf. Empirical Methods Nat. Lang. Process. (EMNLP)*, Hong Kong, China, 2019, pp. 3982–3992, doi: 10.18653/v1/D19-1410. [\[Online\]. Available: https://doi.org/10.18653/v1/D19-1410](https://doi.org/10.18653/v1/D19-1410)
- [31] V. Karpukhin, B. Oğuz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W. Yih, "Dense passage retrieval for open-domain question answering," in *Proc. 2020 Conf. Empirical Methods Nat. Lang. Process. (EMNLP)*, 2020, pp. 6769–6781, doi: 10.18653/v1/2020.emnlp-main.550. [\[Online\]. Available: https://doi.org/10.18653/v1/2020.emnlp-main.550](https://doi.org/10.18653/v1/2020.emnlp-main.550)

- [32] L. Blecher, G. Cucurull, P. P. Fornes, and S. C. M. Montañez, "Nougat: Neural optical understanding for academic documents," in Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR), Seattle, WA, USA, 2024, pp. 15632–15642. [\[Online\]. Available: https://openaccess.thecvf.com/content/CVPR2024/html/Blecher_Nougat_Neural_Optical_Understanding_for_Academic_Documents_CVPR2024_paper.html](https://openaccess.thecvf.com/content/CVPR2024/html/Blecher_Nougat_Neural_Optical_Understanding_for_Academic_Documents_CVPR2024_paper.html)
- [33] D. Wright, D. Wadden, K. Lo, B. Kuehl, A. Cohan, and I. Augenstein, "Generating scientific claims for zero-shot scientific fact checking," in Proc. 60th Annu. Meeting Assoc. Comput. Linguistics (ACL), Dublin, Ireland, 2022, pp. 2448–2460, doi: 10.18653/v1/2022.acl-long.175. [\[Online\]. Available: https://doi.org/10.18653/v1/2022.acl-long.175](https://doi.org/10.18653/v1/2022.acl-long.175)
- [34] J. Wei, X. Wang, D. Schuurmans, M. Bosma, F. Xia, E. Chi, Q. V. Le, and D. Zhou, "Chain-of-thought prompting elicits reasoning in large language models," in Proc. 36th Conf. Neural Inf. Process. Syst. (NeurIPS), New Orleans, LA, USA, 2022, pp. 24824–24837. [\[Online\]. Available: https://proceedings.neurips.cc/paper_files/paper/2022/hash/9d56096ce523622f91bd54cb2506ecac-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2022/hash/9d56096ce523622f91bd54cb2506ecac-Abstract-Conference.html)
- [35] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in Proc. 34th Int. Conf. Mach. Learn. (ICML), Sydney, Australia, 2017, pp. 1321–1330. [\[Online\]. Available: https://proceedings.mlr.press/v70/guo17a.html](https://proceedings.mlr.press/v70/guo17a.html)
- [36] A. N. Angelopoulos and S. Bates, "A gentle introduction to conformal prediction and distribution-free uncertainty quantification," *Foundations and Trends in Machine Learning*, vol. 16, no. 4, pp. 494–591, 2023, doi: 10.1561/22000000094. [\[Online\]. Available: https://doi.org/10.1561/22000000094](https://doi.org/10.1561/22000000094)
- [37] S. Robertson, H. Zaragoza, et al., "The probabilistic relevance framework: BM25 and beyond," *Foundations and Trends in Information Retrieval*, vol. 3, no. 4, pp. 333–389, 2009, doi: 10.1561/15000000019. [\[Online\]. Available: https://doi.org/10.1561/15000000019](https://doi.org/10.1561/15000000019)
- [38] P. He, J. Gao, and W. Chen, "DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing," in Proc. 11th Int. Conf. Learn. Represent. (ICLR), Kigali, Rwanda, 2023. [\[Online\]. Available: https://openreview.net/forum?id=sE7-XhLxHA](https://openreview.net/forum?id=sE7-XhLxHA)
- [39] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. Yih, T. Rocktäschel, S. Riedel, and D. Kiela, "Retrieval-augmented generation for knowledge-intensive NLP tasks," in Proc. 34th Conf. Neural Inf. Process. Syst. (NeurIPS), 2020, pp. 9459–9474. [\[Online\]. Available: https://proceedings.neurips.cc/paper_files/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html](https://proceedings.neurips.cc/paper_files/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html)
- [40] D. S. Himmelstein, A. Lamm, C. Zietz, V. S. Graf, J. A. Doan, and C. S. Greene, "Systematic integration of biomedical knowledge prioritizes drugs for repurposing," *eLife*, vol. 6, p. e26726, 2017, doi: 10.7554/eLife.26726. [\[Online\]. Available: https://doi.org/10.7554/eLife.26726](https://doi.org/10.7554/eLife.26726)
- [41] N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang, "Lost in the middle: How language models use long contexts," *Transactions of the Association for Computational Linguistics*, vol. 12, pp. 157–173, 2024, doi: 10.1162/tacl_a_00638. [\[Online\]. Available: https://doi.org/10.1162/tacl_a_00638](https://doi.org/10.1162/tacl_a_00638)
- [42] M. Sharma, M. Tong, T. Korbak, D. Duvenaud, A. Askill, S. R. Bowman, N. Houlsby, and J. Perez, "Towards understanding sycophancy in language models," in Proc. 2024 Conf. North Amer. Chapter Assoc. Comput. Linguistics: Hum. Lang. Technol. (NAACL), Mexico City, Mexico, 2024, pp. 3120–3139. [\[Online\]. Available: https://aclanthology.org/2024.naacl-long.172/](https://aclanthology.org/2024.naacl-long.172/)
- [43] V. Larivière, S. Haustein, and P. Mongeon, "The oligopoly of academic publishers in the digital era," *PLoS ONE*, vol. 10, no. 6, p. e0127502, 2015, doi: 10.1371/journal.pone.0127502. [\[Online\]. Available: https://doi.org/10.1371/journal.pone.0127502](https://doi.org/10.1371/journal.pone.0127502)
- [44] J. Herzig, P. K. Nowak, T. Müller, F. Piccinno, and J. M. Eisenschlos, "TaPas: Weakly supervised open-domain question answering over tables," in Proc. 58th Annu. Meeting Assoc. Comput. Linguistics (ACL), 2020, pp. 4355–4364, doi: 10.18653/v1/2020.acl-main.398. [\[Online\]. Available: https://doi.org/10.18653/v1/2020.acl-main.398](https://doi.org/10.18653/v1/2020.acl-main.398)
- [45] D. J. Selkoe and J. Hardy, "The amyloid hypothesis of Alzheimer's disease at 25 years," *EMBO Molecular Medicine*, vol. 8, no. 6, pp. 595–608, 2016, doi: 10.15252/emmm.201606210. [\[Online\]. Available: https://doi.org/10.15252/emmm.201606210](https://doi.org/10.15252/emmm.201606210)
- [46] A. Williams, N. Nangia, and S. R. Bowman, "A broad-coverage challenge corpus for sentence understanding through inference," in Proc. 2018 Conf. North Amer. Chapter Assoc. Comput. Linguistics: Hum. Lang. Technol. (NAACL-HLT), New Orleans, LA, USA, 2018, pp. 1112–1122. [\[Online\]. Available: https://aclanthology.org/N18-1101/](https://aclanthology.org/N18-1101/)
- [47] I. Dagan, O. Glickman, and B. Magnini, "The PASCAL recognising textual entailment challenge," in *Machine Learning Challenges. Lecture Notes in Computer Science*, vol. 3944, Berlin, Germany: Springer, 2006, pp. 177–190, doi: 10.1007/11736790_9. [\[Online\]. Available: https://doi.org/10.1007/11736790_9](https://doi.org/10.1007/11736790_9)
- [48] K. Lo, L. L. Wang, M. Neumann, R. Kinney, and D. S. Weld, "S2ORC: The Semantic Scholar Open Research Corpus," in Proc. 58th Annu. Meeting Assoc. Comput. Linguistics (ACL), 2020, pp. 4969–4983, doi: 10.18653/v1/2020.acl-main.447. [\[Online\]. Available: https://doi.org/10.18653/v1/2020.acl-main.447](https://doi.org/10.18653/v1/2020.acl-main.447)
- [49] T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askill, et al., "Language models are few-shot learners," in Proc. 34th Conf. Neural Inf. Process. Syst. (NeurIPS), 2020, pp. 1877–1901. [\[Online\]. Available: https://proceedings.neurips.cc/paper_files/paper/2020/hash/1457c0d6bfc4967418bfb8ac142f64a-Abstract.html](https://proceedings.neurips.cc/paper_files/paper/2020/hash/1457c0d6bfc4967418bfb8ac142f64a-Abstract.html)
- [50] H. W. Harzing, "Are our referencing errors undermining our scholarship and credibility?," *Journal of the American Society for Information Science and Technology*, vol. 53, no. 8, pp. 666–672, 2002, doi: 10.1002/asi.10079. [\[Online\]. Available: https://doi.org/10.1002/asi.10079](https://doi.org/10.1002/asi.10079)